

# Statistical Machine Learning:

## Exercise 3

Yi Cui, 2758172  
Lingwei Liu, 2659255



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

Sommersemester 2020

### Aufgabe 1: Linear Regression

In this exercise, you will implement various kinds of linear regression using the data `lin_reg_train.txt` and `lin_reg_test.txt`. The files contain noisy observations from an unknown function  $f: \mathbb{R} \mapsto \mathbb{R}$ . In both files, the first column represents the inputs and the second column represents the outputs. You can load the data using `numpy.loadtxt`. For all subtasks, assume that the data is identically and independently distributed according to

$$y_i = \Phi(\mathbf{x}_i)^\top \mathbf{w} + \epsilon_i,$$

where

$$\epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

and  $\Phi: \mathbb{R}^1 \rightarrow \mathbb{R}^n$  is a feature transformation such that

$$\mathbf{y} \sim \mathcal{N}(\Phi(\mathbf{X})^\top \mathbf{w}, \sigma^2 \mathbf{I})$$

Additionally, make sure that your implementations support multivariate inputs. The feature transformation are given in each task, if no basis functions are stated explicitly use the data as is  $\Phi(x) = x$ .

1a)

Implement linear ridge regression using linear features, i.e. the data itself. Include an additional input dimension to represent a bias term and use the ridge coefficient  $\mu = 0.01$ .

1. Explain: What is the ridge coefficient and why do we use it? (1)
2. Derive the optimal model parameters by minimizing the squared error loss function. (3)
3. Report the root mean squared error of the train and test data under your linear model with linear features. (2)
4. Include a single plot that shows the training data as black dots and the predicted function as a blue line. (2)

answer:

ridge coefficient lambda in

$$\mathbf{w}_{\text{MAP}} = \left( \Phi \Phi^\top + \frac{\alpha}{\beta} \mathbf{1} \right)^{-1} \Phi \mathbf{y}$$

reason: regularizes the pseudo-inverse

root mean squared error of the training data: 2.9145386993882525

root mean squared error of the training data: 3.8428816992597863

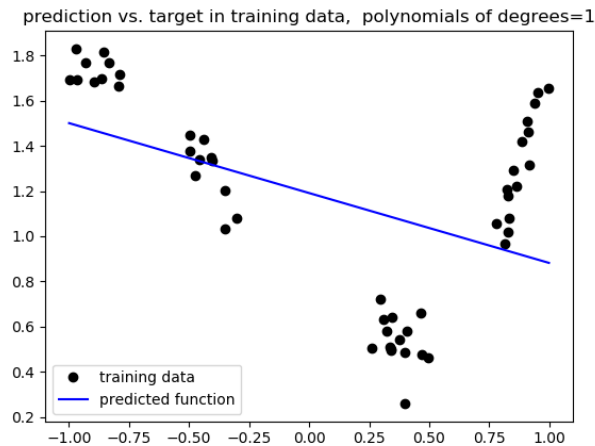


Abbildung 1: Aufgabe1a

1b)

Implement linear ridge regression using linear features, i.e. the data itself. Include an additional input dimension to represent a bias term and use the ridge coefficient  $\mu = 0.01$ .

For polynomials of degrees 2, 3 and 4:

1. Report the root mean squared error of the training data and of the testing data under your model with polynomial features. (2)
2. Include a single plot that shows the training data as black dots and the predicted function as a blue line. (2)
3. Why do we call this method linear regression despite using polynomials? (1)

answer:

polynomials of degrees=1, root mean squared error of the training data: 1.4991687132735394  
 polynomials of degrees=1, root mean squared error of the training data: 2.1687242714148716  
 polynomials of degrees=2, root mean squared error of the training data: 0.6156652380614582  
 polynomials of degrees=2, root mean squared error of the training data: 1.0835803719738029  
 polynomials of degrees=3, root mean squared error of the training data: 0.6152720874799968  
 polynomials of degrees=3, root mean squared error of the training data: 1.0666239820964531

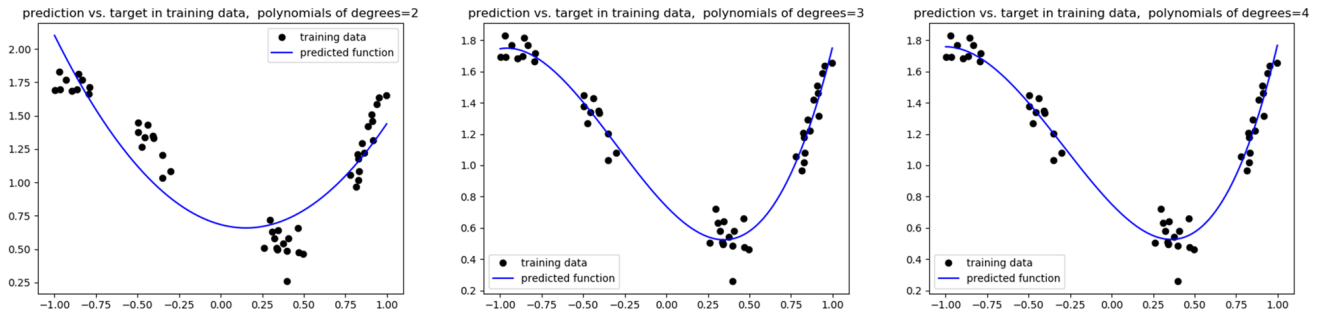


Abbildung 2: Aufgabe1b

and the reason is  $y_{pre} = \Phi_{hub} \cdot w$

1c)

Implement 5-fold cross-validation to select the optimal degree for your polynomial regression.

- Start by splitting the provided data into 5 distinct subsets with each subset consisting of 20% of the original data.
- Use subsets 1 - 4 to train your model with polynomial features of degrees 2, 3 and 4.
- Compute the train RMSEs using your trained models and subsets 1 - 4.
- Compute the validation RMSEs using your trained models and subset 5.
- Compute the test RMSEs using your trained models and the test data.
- Repeat the previous steps, cycling through the subsets until every subset has been used for validation. Provide the following results and answers:

1. For each polynomial degree, report the average train, validation and test RMSEs among all folds. (2)

2. Explain: Do the resulting numbers meet your expectations? Why (not)? (2)

3. Which polynomial degree should be chosen for the given data? Why? (1)

answer:

polynomials of degrees=1, root mean squared error of training: 1.3246142423841534  
 polynomials of degrees=1, root mean squared error of validation: 0.7111489886262025  
 polynomials of degrees=1, root mean squared error of test: 2.183509401119271  
 polynomials of degrees=2, root mean squared error of training: 0.5452281414536173  
 polynomials of degrees=2, root mean squared error of validation: 0.2931779276868292  
 polynomials of degrees=2, root mean squared error of test: 1.092767157004997  
 polynomials of degrees=3, root mean squared error of training: 0.5405756570976916  
 polynomials of degrees=3, root mean squared error of validation: 0.3112176709628346  
 polynomials of degrees=3, root mean squared error of test: 1.0867173876433611

More polynomials of degrees  $\rightarrow$  more precise prediction

Higher order polynomial could have more complex mapping, which can express more details of regression curve. But it could also lead to overfitting.

polynomials of degrees	Mean of RMSE in Test	Std of RMSE in Test
2	2.1835	0.08623
3	1.0928	0.01533
4	1.0867	0.02783

Chose polynomials of degrees 3

Though 4-order polynomials has less Mean of RMSE in cross validation than 3 order, but it shows more Std of RMSE.

1d)

Implement Bayesian linear ridge regression, assuming that  $\mathbf{w}$  follows a multivariate Gaussian distribution, such that

$$\mathbf{w} \sim \mathcal{N}(\boldsymbol{\mu}_0, \boldsymbol{\Lambda}_0^{-1})$$

where ridge regression dictates  $\boldsymbol{\mu}_0 = \mathbf{0}$  and  $\boldsymbol{\Lambda}_0 = \lambda \mathbf{I}$

Here,  $\boldsymbol{\mu}_0$  is the prior weight mean and  $\boldsymbol{\Lambda}_0$  is the prior weight precision matrix, i.e. inverse of covariance matrix. The corresponding posterior parameters can be denoted as  $\boldsymbol{\mu}_n$  and  $\boldsymbol{\Lambda}_n$ .

Assume  $\sigma = 0.1$ , use  $\lambda = 0.01$  and include an additional input dimension to represent a bias term. Use all of the provided training data for a single Bayesian update.

1. State the posterior distribution of the model parameters  $p(\mathbf{w} \mid \mathbf{X}, \mathbf{y})$  (no derivation required). (1)
2. State the predictive distribution  $p(y_* \mid \mathbf{x}_*, \mathbf{X}, \mathbf{y})$  (no derivation required). (1)
3. Report the RMSE of the train and test data under your Bayesian model (use the predictive mean). (1)
4. Report the average log-likelihood of the train and test data under your Bayesian model. (1)
5. Include a single plot that shows the training data as black dots, the mean of the predictive distribution as blue line and 1, 2 and 3 standard deviations of the predictive distribution in shades of blue (you can use matplotlib's `fill_between` function for that). (2)
6. Explain the differences between linear regression and Bayesian linear regression. (1)

answer:

1.

$$p(\mathbf{w} \mid \mathbf{X}, \mathbf{y}, \alpha, \beta) \propto p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}, \beta) p(\mathbf{w} \mid \alpha) \\ \propto p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}, \beta) \mathcal{N}(\mathbf{w} \mid \mathbf{0}, \alpha^{-1} \mathbf{I})$$

where  $p(\mathbf{w} \mid \mathbf{X}, \mathbf{y}, \alpha, \beta)$  is the posterior  $p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}, \beta)$  is the Likelihood of targets under the data and parameters  $p(\mathbf{w} \mid \alpha)$  is the prior

2.

$$p(y_* \mid \mathbf{x}_*, \mathbf{X}, \mathbf{y}) = \int \underbrace{p(y_* \mid \mathbf{x}_*, \theta)}_{\text{likelihood}} \underbrace{p(\theta \mid \mathbf{X}, \mathbf{y})}_{\text{parameter posterior}} d\theta$$

integrate out all possible parameters

RMSE of training data: 2.914538070997487

RMSE of test data: 3.8433605616561968

average log-likelihood of training data: -1.5150672805423742

average log-likelihood of test data: -1.3940163553368101

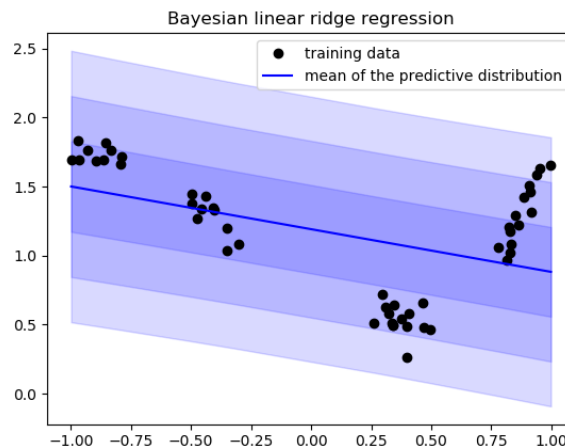


Abbildung 3: Aufgabe1d

1e)

Implement Bayesian linear ridge regression using squared exponential (SE) features. In other words, replace your observed data matrix  $X \in \mathbb{R}^{n \times 1}$ , where:

$$\Phi_{ij} = \exp \left( -\frac{1}{2} \beta (X_i - \alpha_j)^2 \right).$$

Set  $k = 20$ ,  $\alpha_j = j * 0.1 - 1$  and

$\beta = 10$ . Use the ridge coefficient  $\lambda = 0.01$  and assume known Gaussian noise with  $\sigma = 0.1$ . Include an additional input dimension to represent a bias term.

1. Report the RMSE of the train and test data under your Bayesian model with SE features. (1)

2. Report the average log-likelihood of the train and test data under your Bayesian model with SE features. (1)

3. Include a single plot that shows the training data as black dots, the mean of the predictive distribution as blue line and 1, 2 and 3 standard deviations of the predictive distribution in shades of blue (you can use matplotlib's `fill_between` function for that). (2)

4. How can SE features be interpreted from a statisticians point of view? What are  $\alpha$  and  $\beta$  in that context? (2)

answer:

RMSE of training data: 0.5827509166240352

RMSE of test data: 1.1618310002287155

average log-likelihood of training data: -0.10392644293632658

average log-likelihood of test data: -0.22894630923741704

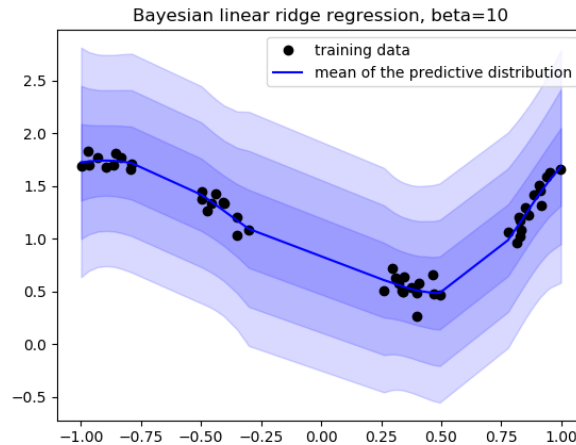


Abbildung 4: Aufgabe1e

1f)

In this bonus assignment, you will perform a grid search using

$\beta \in \{1, 10, 100\}$  to select a 'better'  $\beta$  for your squared exponential features from the previous subtask. Keep using the same settings as in the previous subtask, except  $\beta$

. Grid search is a simple method to select hyperparameters, such as  $\beta$

. First, a list of possible values, or a grid, in case of multiple hyperparameters, is created to define a discrete search space. For every value in this search space, a model is trained and evaluated using a score or loss function. Finally, the hyperparameters that yield the highest score or smallest loss are selected. Here, the log-marginal likelihood will be used as a score function.

The log-marginal likelihood of our Bayesian linear model can be expressed as (see Bishop Ch. 3.5)

$$\begin{aligned} \log p(\mathbf{y} \mid \mathbf{X}) &= \log \int p(\mathbf{y} \mid \mathbf{X}, \mathbf{w}) p(\mathbf{w}) d\mathbf{w}, \\ &= \frac{k+1}{2} \log \lambda - \frac{n}{2} \log \sigma^2 - \frac{1}{2} \frac{\|\mathbf{y} - \Phi \boldsymbol{\mu}\|_2^2}{\sigma^2} + \frac{\lambda}{2} \boldsymbol{\mu}^\top \boldsymbol{\mu} - \frac{1}{2} \log |\boldsymbol{\Lambda}| - \frac{n}{2} \log 2\pi \end{aligned}$$

where

$$\begin{aligned} \boldsymbol{\mu} &= \sigma^{-2} \boldsymbol{\Lambda}^{-1} \Phi^\top \mathbf{y} \\ \boldsymbol{\Lambda} &= \sigma^{-2} \Phi^\top \Phi + \lambda \mathbf{I} \end{aligned}$$

Here,  $k$  is the dimensionality of the feature space, the  $+1$  comes from the extra bias term. Using the marginal likelihood to select hyperparameters is typically referred to as empirical Bayes, type-II maximum likelihood or evidence approximation. Applying the logarithm (analytically) ensures numerical stability.

1. What is the difference between the marginal likelihood  $p(y|X)$  and the likelihood  $p(y|X, w)$ ? (1)

2. For each  $\beta$

, report RMSE and average log-likelihood of the train and test data and the log-marginal likelihood. (3)

3. For each beta, include a single plot that shows the training data as black dots, the mean of the predictive distribution as blue line and 1, 2 and 3 standard deviations of the predictive distribution in shades of blue (you can use matplotlib's `fill_between` function for that). (2)

4. According to the grid search, which value for  $\beta$

is the best? Why? (2)

5. Compare the log-marginal likelihood values to the average train and test log-likelihood values. What do you observe? Is the log-marginal likelihood a 'good' score function compared to the train log-likelihood? (2)

answer:

SE	1	10	100
RMSE of training data	0.9174	0.5828	0.5498
average log-likelihood of training data	-0.2917	-0.1039	-0.0856
RMSE of test data	1.5690	1.1618	1.6271
average log-likelihood of test data	-0.3976	-0.2289	-0.3902
log-marginal likelihood of training data	-22.3919	1.1271	3.3718
log-marginal likelihood of test data	-103.3937	-49.3719	-112.0934

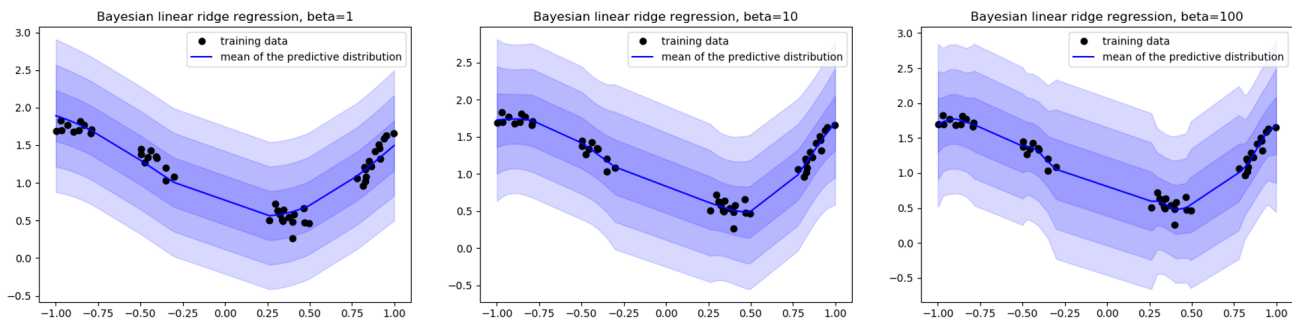


Abbildung 5: Aufgabe1f

according to RMSE in Test, best beta is 100

according to log-likelihood in Test, best beta is 10

according to log-marginal likelihood in Test, best beta is 10

Although beta = 100 performs best in RMSE, its variance is significantly larger. Based on the results of log-likelihood and log-marginal likelihood, the best beta should be 10

Log-marginal likelihood includes consideration of the weight matrix  $\mu$  and ridge coefficient.

In terms of loss function (ie, score function), the occurrence of overfitting can be reduced to a certain extent.

---

### Aufgabe 2: Linear Classification

---

In this exercise, you will use the dataset `ldaData.txt`, containing 137 feature points  $\vec{x}$ . The first 50 points belong to class  $C_1$ , the second 43 to class  $C_2$ , the last 44 to class  $C_3$ .

---

2a)

---

Explain the difference between discriminative and generative models and give an example for each case. Which model category is generally easier to learn and why?

Answer:

In General, A Discriminative model models the decision boundary between the classes. A Generative Model explicitly models the actual distribution of each class.

A Generative Model learns the joint probability distribution  $p(x, y)$ . It predicts the conditional probability with the help of Bayes Theorem. A Discriminative model learns the conditional probability distribution  $p(y|x)$ . Both of these models were generally used in supervised learning problems.

In general generative models are easier to learn than discriminative models. Because in practice, it is usually hard to compute the posterior for discriminative models.



2b)

Use Linear Discriminant Analysis to classify the points in the dataset. Attach two plots with the data points using a different color for each class: one plot with the original dataset, one with the samples classified according to your LDA classifier. Attach a snippet of your code and discuss the results. How many samples are misclassified? (You are allowed to use built-in functions for computing the mean and the covariance.)

answer:

# -\*- coding: utf-8 -\*-

```
import numpy as np
import matplotlib.pyplot as plt
```

```
def load_data(filename):
    data = np.loadtxt(f"D:\\Test\\dataSets\\{filename}.txt")
    C_1 = np.mat(data[:50, :]).T
    C_2 = np.mat(data[50:93, :]).T
    C_3 = np.mat(data[93:, :]).T
    return C_1, C_2, C_3, data
```

```
def normal_vec_w (c_1, c_2):
    mean_c_1 = np.mean(c_1, 1)
    mean_c_2 = np.mean(c_2, 1)
    diff_w = np.dot ((c_1 - mean_c_1), ((c_1 - mean_c_1).T)) + np.dot ((c_2 - mean_c_2), ((c_2 - mean_c_2).T))
    mat_w = np.dot(np.linalg.inv(diff_w), (mean_c_2 - mean_c_1))
    normal_w = mat_w[0] / mat_w[1]
    return mat_w, normal_w
```

```
def offset(w, c_1, c_2):
    c1_re = np.dot(w.T, c_1)
    c2_re = np.dot(w.T, c_2)
    mean_c1_re = np.mean(c1_re)
    mean_c2_re = np.mean(c2_re)
    var_c1_re = np.var(c1_re)
    var_c2_re = np.var(c2_re)
    r1 = 1 / (2 * var_c1_re) - 1 / (2 * var_c2_re)
    r2 = (mean_c2_re / var_c2_re) - (mean_c1_re / var_c1_re)
    r3 = mean_c1_re ** 2 / (2 * var_c1_re) - mean_c2_re ** 2 / (2 * var_c2_re) -
        np.log(np.sqrt(var_c2_re/var_c1_re))

    root_r = np.roots([r1, r2, r3])
    for i in range (len(root_r)):
        if root_r[i] < max(mean_c1_re, mean_c2_re) and root_r[i] > min(mean_c1_re, mean_c2_re):
            w0 = root_r[i]
            break
    return w0
```

```
if __name__ == "__main__":
    C_1, C_2, C_3, data = load_data('ldaData')
    a = C_1.T
    b = C_2.T
    c = C_3.T

    mat_w12, normal_w12 = normal_vec_w(C_1, C_2)
    mat_w13, normal_w13 = normal_vec_w(C_1, C_3)
```

```

mat_w23, normal_w23 = normal_vec_w(C_2,C_3)

w0_12 = offset(mat_w12,C_1, C_2)
w0_13 = offset(mat_w13,C_1, C_3)
w0_23 = offset(mat_w23,C_2, C_3)

#without LDF
for i in range(len(a)):
    plt.scatter(a[i, 0], a[i, 1] , marker = 'v' , color = 'red')
    i = i + 1

for i in range(len(b)):
    plt.scatter(b[i, 0], b[i, 1] , marker = 'x' , color = 'yellow')
    i = i + 1

for i in range(len(c)):
    plt.scatter(c[i, 0], c[i, 1] , marker = 'o' , color = 'blue')
    i = i + 1

plt.show()

# with LDF
for i in range(len( data )):
    if np.dot(mat_w12.T, np.mat(data[i, :]).T) - w0_12 < 0:
        plt.scatter(data[i, 0], data[i, 1] , marker = 'v' , color = 'red')
    elif np.dot(mat_w23.T, np.mat(data[i, :]).T) - w0_23 < 0:
        plt.scatter(data[i, 0], data[i, 1], marker = 'x', color = 'yellow')
    else:
        plt.scatter(data[i, 0], data[i, 1], marker = 'o', color = 'blue')
    i = i + 1
plt.show()

```

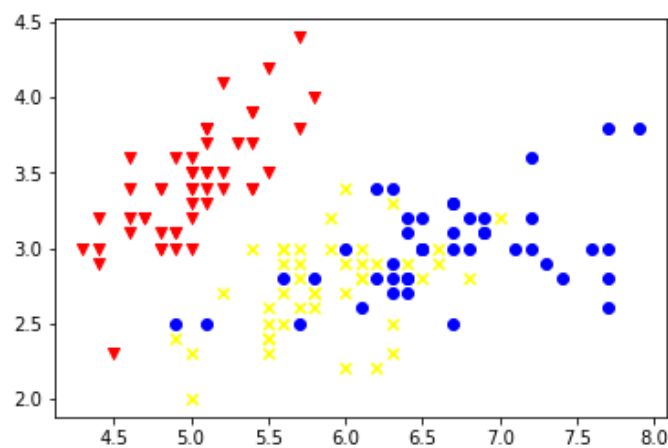


Abbildung 6: Aufgabe2b\_original\_classes

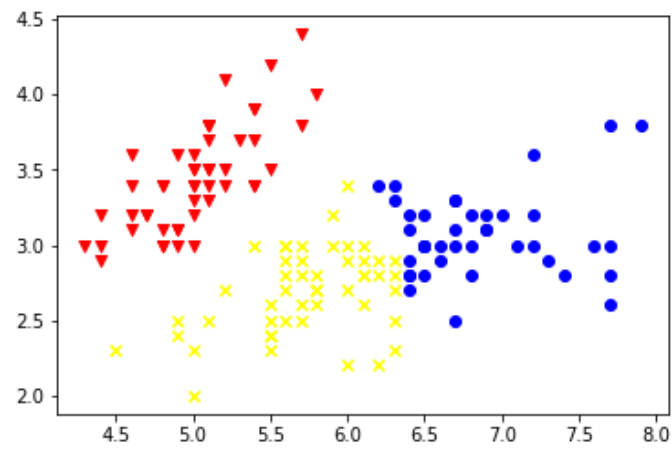


Abbildung 7: Aufgabe2b\_classes\_with\_LDA

**Aufgabe 3: Principal Component Analysis**

In this exercise, you will use the dataset iris.txt. It contains data from three kind of Iris flowers ('Setosa', 'Versicolour' and 'Virginica') with 4 attributes: sepal length, sepal width, petal length, and petal width. Each row contains a sample while the last attribute is the label (0 means that the sample comes from a 'Setosa' plant, 1 from a 'Versicolour' and 2 from 'Virginica'). (You are allowed to use built-in functions for computing the mean, the covariance, eigenvalues, eigenvectors and singular value decomposition.)

3a)

Normalizing the data is a common practice in machine learning. Normalize the provided dataset such that it has zero mean and unit variance per dimension. Why is normalizing important? Attach a snippet of your code.

answer:

Normalization is important in PCA since it is a variance maximizing exercise. It projects the original data onto directions which maximize the variance.

And it will give more emphasis to those variables having higher variances than to those variables with very low variances while identifying the right principle component.

```
def normalization(dataset):
    x , y = np.shape(dataset)
    for i in range( y - 1):
        dataset[:, i] = dataset[:, i] - np.mean(dataset[:, i])
        dataset[:, i] = dataset[:, i] * np.sqrt(x / float(np.dot(np.mat(dataset[:, i]),
            np.mat(dataset[:, i]).T)))
    return x, y, dataset
```

3b)

Apply PCA on your normalized dataset and generate a plot showing the proportion (percentage) of the cumulative variance explained. How many components do you need in order to explain at least 95% of the dataset variance? Attach a snippet of your code.

answer:

```
# -*- coding: utf-8 -*-

import numpy as np
import matplotlib.pyplot as plt

def load_data(filename):
    data = np.loadtxt(f"D:\\Test\\dataSets\\{filename}.txt" , delimiter = ',')
    return data

def normalization(dataset):
    x , y = np.shape(dataset)
    for i in range( y - 1):
        dataset[:, i] = dataset[:, i] - np.mean(dataset[:, i])
        dataset[:, i] = dataset[:, i] * np.sqrt(x / float(np.dot(np.mat(dataset[:, i]),
            np.mat(dataset[:, i]).T)))
    return x, y, dataset
```

---

```

if __name__ == "__main__":
    data_raw = load_data('iris')
    N, M, data_norm = normalization(data_raw)

    data_cal = np.mat(data_norm[:, 0 : 4]).T
    cov_data = np.cov(data_cal)
    lambda_data, W = np.linalg.eig(cov_data)

    #3b
    lambda_norm = lambda_data / np.sum(lambda_data)
    a = np.linspace(1, M-1, M-1)
    b = np.zeros(M-1)
    for i in range(M-1):
        b[i] = np.sum(lambda_norm[: i+1])
    plt.plot(a, b)
    plt.show()

```

---

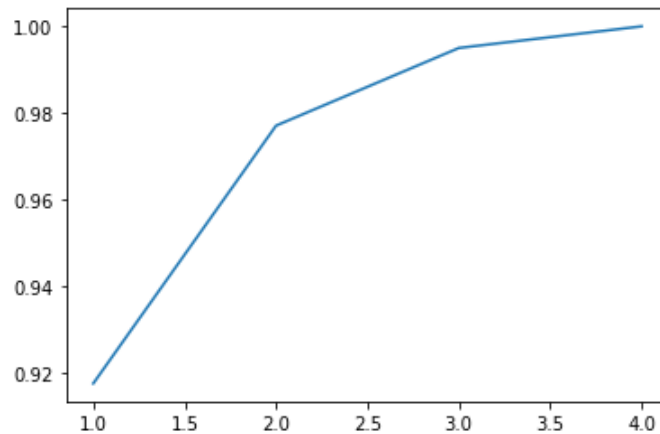


Abbildung 8: Aufgabe3b

the two first varicance are needed, to explain at least 95% of the dataset variance, as it is seen in the picture.

---

3c)

---

Using as many components as needed to explain 95% of the dataset variance, generate a scatter plot of the lowerdimensional projection of the data. Use different colors or symbols for data points from different classes. What do you observe? Attach a snippet of your code.

answer:

---

```

# -*- coding: utf-8 -*-

import numpy as np
import matplotlib.pyplot as plt

def load_data(filename):
    data = np.loadtxt(f"D:\\Test\\dataSets\\{filename}.txt" , delimiter = ',')

```

```

    return data

def normalization(dataset):
    x, y = np.shape(dataset)
    for i in range(y - 1):
        dataset[:, i] = dataset[:, i] - np.mean(dataset[:, i])
        dataset[:, i] = dataset[:, i] * np.sqrt(x / float(np.dot(np.mat(dataset[:, i]),
            np.mat(dataset[:, i]).T)))
    return x, y, dataset

if __name__ == "__main__":
    data_raw = load_data('iris')
    N, M, data_norm = normalization(data_raw)

    data_cal = np.mat(data_norm[:, 0 : 4]).T
    cov_data = np.cov(data_cal)
    lambda_data, W = np.linalg.eig(cov_data)

    #3c
    B = np.mat(W[:, 0: 2])
    class_2d = np.dot(B.T, (data_cal - np.mean(data_cal, 1)))
    for i in range(N):
        if data_norm[i, 4] == 0:
            plt.scatter(class_2d[0, i], class_2d[1, i], marker = 'v', color = 'red')
        elif data_norm[i, 4] == 1:
            plt.scatter(class_2d[0, i], class_2d[1, i], marker = 'x', color = 'yellow')
        else:
            plt.scatter(class_2d[0, i], class_2d[1, i], marker = 'o', color = 'blue')
    plt.show()

```

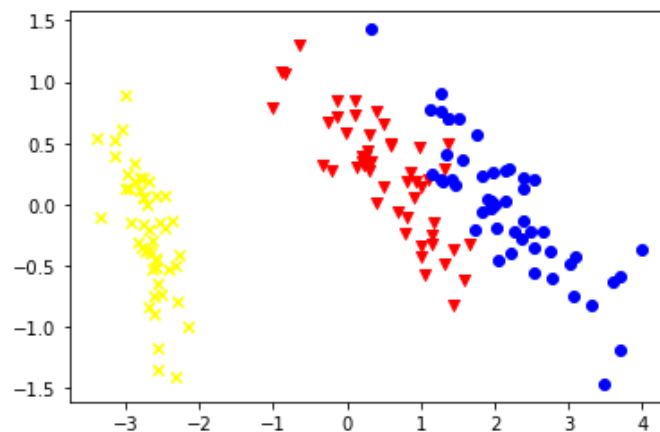


Abbildung 9: Aufgabe3c

3d)

Reconstruct the original dataset by using different number of principal components. Using the normalized root mean square error (NRMSE) as a metric, fill the table below (error per input versus the amount of principal components used).

answer:

N. of components	$x_1$	$x_2$	$x_3$	$x_4$
1				
2				
3				
4				

Attach a snippet of your code. (Remember that in the first step you normalized the data.)

answer:

```
# -*- coding: utf-8 -*-
```

```
import numpy as np
import matplotlib.pyplot as plt
```

```
def load_data(filename):
    data = np.loadtxt(f"D:\\Test\\dataSets\\{filename}.txt" , delimiter = ',')
    return data

if __name__ == "__main__":
    data_raw = load_data('iris')
    #3d
    data_cal_raw = np.mat(data_raw[:, 0 : 4]).T
    error = np.mat(np.zeros((4,4)))
    cov_data_raw = np.cov(data_cal_raw)
    lambda_data_raw , W_raw = np.linalg.eig(cov_data_raw)
    for i in range(M - 1):
        B_raw = np.mat(W_raw[:, 0: i+1])
        class_2d_raw = np.dot(B_raw.T, (data_cal_raw - np.mean(data_cal_raw, 1)))
        datau_cal = np.mean(data_cal_raw, 1) + np.dot(B_raw , class_2d_raw)
        error[i, :] = np.sqrt(np.sum(np.power((data_cal_raw - datau_cal), 2), 1) / N).T
    print(error)
```

As a result the table will be filled:

N. of components	$x_1$	$x_2$	$x_3$	$x_4$
1	0.4137	0.4010	0.1615	0.2058
2	0.1485	0.2131	0.0813	0.1929
3	0.0389	0.0496	0.0753	0.1209
4	0	0	0	0

3e)

In machine learning, it is often desirable to ‘whiten’ the data before applying a model or an algorithm. In this context, ‘whitening’ the data refers to a transformation that warps the data into a spherical shape, such that the data dimensions become uncorrelated and the individual means and variances are 0 and 1, respectively. In particular, PCA can be used to compute such a transformation.

Recommended reading: [ufldl.stanford.edu/tutorial/unsupervised/PCAWWhitening](http://ufldl.stanford.edu/tutorial/unsupervised/PCAWWhitening)

1. Explain the difference between PCA and ZCA whitening. (1)
2. State the equation(s) to compute the ZCA whitening parameters, given the data. (1)

3. State the equation(s) to whiten a (new) data example  $x$ , given the ZCA parameters. (1)

4. Compute and report the ZCA whitening parameters for the unnormalized IRIS data (including numerical values!). For numerical stability, use  $\epsilon = 1e - 5$  (2)

answer:

1. For PCA whitening, the Covariance is exact 1, and for ZCA whitening, the Covariance is same but not strict 1. And PCA whitening can be done to reduce the dimension, while ZCA is more used to uncorrelate.

2.

$$X_{ZCAwhitening} = U X_{PCAWhite,i}$$

3.

$$\begin{aligned}\Sigma &= \frac{1}{m} X X^T \\ \frac{1}{m} X X^T &= U A U^T \\ X_{rot} &= U^T X \\ X_{PCAwhite,i} &= \frac{X_{rot,i}}{\sqrt{\lambda_i + \epsilon}}, \epsilon = 1.0e^{-5} \\ X_{ZCAwhitening} &= U X_{PCAWhite,i}\end{aligned}$$

3f)

Throughout this class we have seen that PCA is an easy and efficient way to reduce the dimensionality of some data. However, it is able to detect only linear dependences among data points. A more sophisticated extension to PCA, Kernel PCA, is able to overcome this limitation. This question asks you to deepen this topic by conducting some research by yourself: explain what Kernel PCA is, how it works and what are its main limitations. Be as concise (but clear) as possible.

answer:

Kernel PCA: an extension of principal component analysis (PCA) using techniques of kernel methods. Using a kernel, the originally linear operations of PCA are performed in a reproducing kernel Hilbert space.

Process: Solve the following eigenvalue problem:

$$K u_i = \lambda_i u_i, \lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_N$$

The projection of the test sample  $\Phi(X_j)$  on the  $i$ -th eigenvector can be computed by

$$K u_i = \lambda_i u_i, \lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_N v_i^T \Phi(x_j) = \frac{1}{\sqrt{\lambda_j}} u_i^T \begin{bmatrix} K(x_1, x_j) \\ \vdots \\ K(x_N, x_j) \end{bmatrix}$$

limitations:

- a) PCA assumes a linear transformation:  $\rightarrow$  With centering of data, one can only do a rotation in space.
- b) It fails at finding directions that require a non-linear transformation.